

Divi Jaiwanth

+91 8374802689 | divijaiwanth@gmail.com | linkedin.com/divijaiwanth | github.com/divijaiwanth

EDUCATION

Manipal Institute of Technology

B.Tech in Computer Science and Engineering | CGPA: 8.29/10.00

Bengaluru, India

June 2023 – July 2027

Indian Institute of Technology Madras (Online)

B.S. in Data Science and Programming

Chennai, India

May 2023 – July 2027

EXPERIENCE

ML Engineer Intern

Rewardsy

May 2025 – July 2025

Bengaluru, Karnataka

- Built a custom QR code generator using Stable Diffusion ControlNet and PyTorch, improving QR design quality and enabling 30% faster generation time.
- Deployed the application using FastAPI and integrated Firebase for cloud storage, reducing data retrieval time by 25% and lowering server maintenance efforts.
- Implemented error handling and memory-efficient processing, resulting in 40% fewer system failures and increased reliability during peak operations.

PROJECTS

Lumora - Zero-Shot Event Face Recognition Pipeline | [RetinaFace](#), [ArcFace](#), [FAISS](#), [OpenCV](#), [Streamlit](#) | [GitHub](#)

- Engineered a production-grade zero-shot face recognition system for dense event photography using RetinaFace detection, 5-point affine alignment, and 512-dim ArcFace embeddings - no training data required per event.
- Indexed embeddings in FAISS (IndexFlatIP) for sub-second cosine similarity search across thousands of event photos.
- Shipped a Streamlit app enabling selfie-based photo retrieval with bounding box overlays on full-resolution event images.

Marot - Autonomous Discord Marketing Agent | [Python](#), [Mistral](#), [Ollama](#), [Mem0](#), [RAG](#), [discord.py](#) | [GitHub](#)

- Architected a dual-layer AI agent separating persistent user memory (Mem0) from a keyword-optimized RAG knowledge base, maintaining domain accuracy and eliminating hallucinations in live Discord conversations.
- Replaced LangChain with a custom async orchestrator over Ollama's Python API, cutting inference latency and removing framework overhead.
- Deployed fully local LLM stack (Mistral via Ollama) with zero external API calls, achieving cost-free inference at production scale.

RAGTube - YouTube RAG System | [Python](#), [LangChain](#), [FAISS](#), [HuggingFace](#), [Ollama](#), [Streamlit](#) | [GitHub](#)

- Built an end-to-end RAG pipeline that ingests YouTube transcripts, chunks and embeds them via HuggingFace sentence-transformers, and stores vectors in FAISS for semantic Q&A over long-form video content.
- Ran fully local LLM inference via Ollama (Mistral) with zero API dependency, demonstrating practical experience with embedding models, vector stores, and retrieval-augmented generation.
- Delivered a Streamlit interface for URL-based ingestion and natural language querying, making the system accessible without any CLI usage.

TECHNICAL SKILLS

Languages: Python, Java, SQL, JavaScript

Cloud & Infrastructure: AWS S3, AWS Bedrock, AWS Lambda, Firebase, REST APIs, CI/CD

Frameworks & Tools: FastAPI, LangChain, Streamlit, Node.js, Git, FAISS, Ollama, Vector Databases

Libraries: PyTorch, OpenCV, HuggingFace Transformers, RetinaFace, ArcFace, Keras, TensorFlow

AI/ML: Computer Vision, Face Recognition, Generative AI, RAG, Stable Diffusion, LLMs

CERTIFICATIONS

Build Basic Generative Adversarial Networks (GANs) | [DeepLearning.AI](#), 2024 | [Certificate](#)

Generative AI Fundamentals | [Google](#), 2023 | [Certificate](#)

Machine Learning with Python | [IBM](#), 2023 | [Certificate](#)